IASA Masterclass · Canvas guide

IASA canvases, explained with AIForge

What each canvas is for — and the same canvas filled in for our own AI developer.

First
how canvases help

Then
14 canvases, one by one

Each time
the AIForge example

Roles only — no names, no customer detail.

How canvases help

Six reasons to fill a one-page canvas instead of writing a long document

Start here
Why bother?

1 One page, not a document

A single page gets read and kept up to date. A fifty-page document does neither.

2 A picture beats a paragraph

Words alone are ambiguous. A box you must fill forces a clear answer.

3 Everyone works on one wall

Print it, pin it up, fill it with sticky notes. The whole team sees the same thing.

4 Each canvas feeds the next

The persona drives the journey; the journey drives the requirements. Cause and effect.

5 Decisions stay traceable

Requirement, decision, measure — linked. Nobody re-argues the choice next year.

6 Gaps show up early

An empty box is a question nobody has answered yet. Better found now than in the build.

In short: a canvas is a one-page question sheet. Fill the boxes and the argument can be checked.

The canvases at a glance

14 canvases, grouped by the question each one answers

The map
Five questions

Why?
the business

1. **Business Model Canvas**

Who?
the people

2. **Persona Card**
3. **Customer Journey Map**
4. **Power / Interest Grid**

What?
the service

5. **Use Case Card**
6. **Context View**
7. **Service Blueprint**
8. **Capability Map**
9. **Capability Card**

How?
the design

10. **ASR Card**
11. **QATT Card**
12. **Decision Record (ADR)**

Where next?
the plan

13. **OKR Card**
14. **Roadmap Canvas**

In short: start with why and who. The design canvases only make sense after those.

Our example: AIForge

Running example
AIForge

One input, one output — and nobody has to watch in between

A ticket

Written in plain words, the way you would ask a colleague

AIForge

A small team of AI agents on our own GPUs plans, writes, tests and reviews the change

A pull request

A tested, reviewed code change, ready for a person to read and merge

Our machines only

No cloud. No code or data leaves the building.

Seven small agents

Each has one small job and only the context it needs.

A person acts twice

Write the ticket. Read the pull request.

In short: keep this picture in mind — every canvas that follows describes it from a different side.

Business Model Canvas

Why?
Canvas 1 of 14

It answers: how does this create value, and for whom?

WHAT IT IS

Nine boxes that show how a product or service earns its keep.

HOW IT HELPS

- The whole business case on one page
- Shows who benefits and what it costs
- Starts the chain: segments become personas

KEY PARTNERS

IT operations
QA and testers

KEY ACTIVITIES

Plan, write, test, review

KEY RESOURCES

Our own GPUs
Local models
Run memory

VALUE OFFERED

Ticket in, tested pull request out.

Code never leaves the building.

RELATIONSHIP

An honest verdict on every pull request

CHANNELS

Chat
Ticket tracker

WHO IT SERVES

Delivery engineers
Junior developers
Engineering leads

WHAT IT COSTS

GPUs, model calls, and one operator to run the platform

WHAT COMES BACK

Senior hours freed for the hard work; more delivered by the same team

In short: AIForge sells nothing. It gives senior hours back, inside our own walls.

Persona Card

Who?
Canvas 2 of 14

It answers: who exactly are we building for?

WHAT IT IS

A one-page picture of one real kind of user, in their own words.

HOW IT HELPS

- Stops “the user” meaning everyone
- Their objection is part of the person
- Drives the journey and the requirements



PRIMARY PERSONA

Delivery engineer

“Half my week goes on work I already know how to do.”

THEIR DAY

- Owns a slice of a live system
- Three to six small tickets a week
- Interrupted often; deep work is rare

WHAT THEY WANT

- Boilerplate and glue done for them
- Keep the interesting design work
- Never hand code to an outside service

WHAT BLOCKS THEM

- A pull request they do not trust
- A tool that half-finishes and stays quiet
- Setup work before the tool is useful

WHAT EARNS TRUST

- It compiles, or it says it failed
- “I could not do this” in the first line
- A second model checked the work

In short: it is built for the delivery engineer. If their day does not change, nothing else matters.

Customer Journey Map

Who?
Canvas 3 of 14

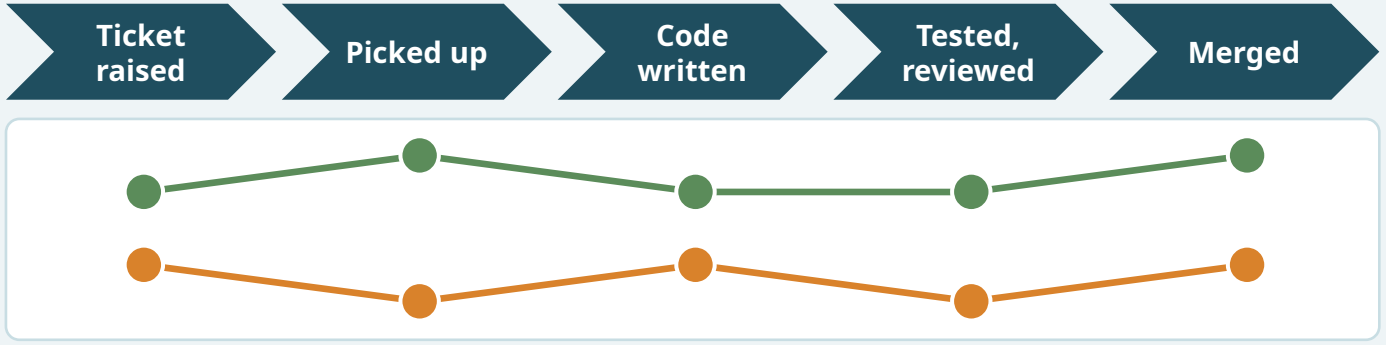
It answers: where does the person wait, and how does it feel?

WHAT IT IS

The steps a person goes through, with how each step feels.

HOW IT HELPS

- Finds the wait, not just the step
- Before and after sit side by side
- Shows the moment that matters



Today	sits in the backlog	waits for a senior to be free	squeezed between other work	rushed at the end	days after it was raised
With AIForge	picked up as written	a run starts, nobody waits	agents write it in a copy of the code	tests first, second model reviews	pull request ready to read

In short: the cost was never the typing — it was the waiting. AIForge removes the wait.

Power / Interest Grid

Who?
Canvas 4 of 14

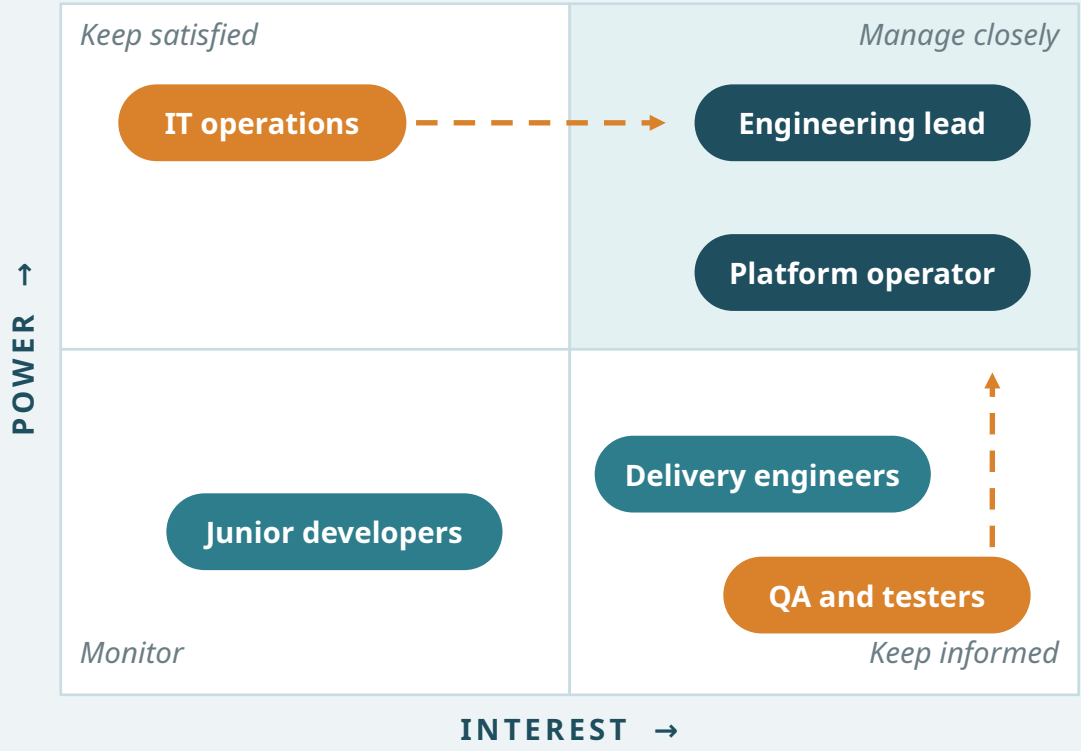
It answers: who must we win over, and who must own it?

WHAT IT IS

A two-by-two: each group placed by its say and its stake.

HOW IT HELPS

- Shows where each group stands today
- Arrows show where they need to be
- Turns “people issues” into named moves



The moves we need

IT operations
watching → owning the model service

QA and testers
affected → owning the test gate

Platform operator
one person → two

In short: the tool works. What is left is people — three jobs still have no owner.

Use Case Card

What?
Canvas 5 of 14

It answers: who starts it, what happens, and what comes out?

WHAT IT IS

One job the system does, told as actor, trigger, steps and outcome.

HOW IT HELPS

- Names the actor and the trigger
- Says what must be true before it starts
- Ends in an outcome you can check

ACTOR AND TRIGGER



Delivery engineer
files a ticket in the tracker

Runs several times a week

MUST ALREADY BE TRUE

- The ticket names a real project
- The models are loaded and warm
- The repo builds before we start

MAIN FLOW

1. Someone files a ticket in plain words
2. The ticket is rewritten so it is clear
3. The work is split; a shared spec is written
4. Each part is built in its own copy of the repo
5. Tests are written, then run against the build
6. A second model reviews the change
7. A pull request opens with an honest verdict

OUTCOME

A pull request
a person reviews

Or a first line
saying it could not

Either way
a run history for audit

In short: the senior spends minutes reading instead of hours writing.

Context View

What?
Canvas 6 of 14

It answers: what is inside the system, and what does it touch?

WHAT IT IS

The system as one box, with everything that goes in and comes out.

HOW IT HELPS

- Draws the boundary everyone argues about
- Lists assumptions next to the boxes
- Makes the risks at the edges visible

GOES IN

- Tickets from the tracker
- The repo and its history
- Past runs and lessons

our own hardware

AIForge
Plans the work
Writes and tests code
Reviews and merges

COMES OUT

- A pull request
- A run history for audit
- New lessons stored

WHAT WE ASSUME

- Models are loaded the right way
- Models stay warm, or runs stop mid-way
- The ticket names a real project

WHAT TO BEWARE OF

- About 3–5 files is today's ceiling
- Some languages still beat the smaller models
- Shared GPUs limit how many tickets run at once

In short: write the unknowns down, not just the boxes — the edges are where runs fail.

Service Blueprint

What?
Canvas 7 of 14

It answers: what does the person see, and what runs behind it?

WHAT IT IS

The service drawn in lanes: what the person sees on top, the machinery below.

HOW IT HELPS

- The line of visibility is the point
- Shows who does what at each stage
- Finds work nobody owns

	1. Ask	2. Work	3. Check	4. Hand back
The person sees	the ticket they wrote	a run in progress	test output	a pull request and a verdict
The person does	files it, in plain words	nothing — goes elsewhere	nothing	reads, comments, merges
Front stage	ticket is picked up	branch and worktree open	compile gate runs	pull request opens
line of visibility — below this, nobody is watching				
Back stage: agents	enhance and clarify	plan, spec, write	test, review, reconcile	write the verdict
Platform	ticket store and routing	model service on our GPUs	build tools, test runners	run history, lessons kept

In short: above the line there are two moments: file it, read it. The design work is all below.

Business Capability Map

What?
Canvas 8 of 14

It answers: what must the business be good at — and who does it today?

WHAT IT IS

A grid of everything the business must be able to do. What, never how.

HOW IT HELPS

- One view of the whole job, end to end
- Colour shows where help already lands
- The uncoloured boxes are the plan

Understand demand	Design the change	Build the change	Prove quality	Release and operate	Learn and improve
Capture a ticket	Read the codebase	Write the code	Write tests	Open a pull request	Keep lessons
Clarify what is wanted	Split the work	Wire files together	Run the build	Merge	Share across the team
Prioritise	Agree the contract	Keep to standards	Second-pair review	Run the platform	Measure adoption

AIForge does this today · Partly — being built · Still all people

In short: 11 of 18 done by AIForge today, 3 partly. The rest is an ownership gap.

Capability Card

What?
Canvas 9 of 14

It answers: what does one capability need — people, process, technology?

WHAT IT IS

One capability taken apart: who supports it, how it runs, how mature it is.

HOW IT HELPS

- People, process, technology — in that order
- A maturity level you can argue about
- The gap to target becomes the to-do list

BUILD THE CHANGE

Turn an agreed spec into code that compiles and fits the repo it lands in.

MATURITY

Growing | **Good** | Strong

1 People

- One platform operator today
- Delivery engineer reviews the output
- Nobody named for the model service

2 Process

- Spec and contract written before code
- Each part built in its own worktree
- Only compiling code reaches a merge

3 Technology

- Local models on hardware we own
- Repo map and past-run memory
- Build and test tools already in use

GAP TO TARGET

Ceiling is 3–5 files; target is 8–10

Memory not yet shared across the team

Second operator not trained

In short: raise it with bigger local models, shared memory, and a named owner for the platform.

ASR Card

How?
Canvas 10 of 14

It answers: which few requirements shape the whole design?

WHAT IT IS

ASR = architecturally significant requirement: a rule the design cannot break.

HOW IT HELPS

- Picks the few rules that really matter
- States the risk next to the rule
- Each one later needs a decision and a measure

1 HIGH

Do the job on its own

Ticket in, pull request out. Nobody in the middle.

RISK Bad automatic pull requests destroy trust fast.

PROOF
The first full run went from ticket to pull request with no help.

2 HIGH

Make the files fit together

Change five files and they still build as one piece.

RISK Without it, big tickets look done and do not compile.

PROOF
0 working linked lines → 53, once a shared spec was added.

3 MEDIUM-HIGH

Stay on our machines

Hardware we own. No cloud. Works with the internet off.

COST Local models are weaker than hosted ones.

PROOF
Zero runs have sent data outside our network.

In short: three rules shaped every design choice. Break any one and the tool is not worth using.

QATT Card

How?
Canvas 11 of 14

It answers: how do we turn “it must be good” into a number?

WHAT IT IS

A quality written as a test scenario: who does what, and the measured response.

HOW IT HELPS

- Turns a wish into a pass or fail
- Six fixed parts, so nothing is vague
- Gives every ASR a way to be checked

SCENARIO **A multi-file ticket comes out as one working change**

SOURCE	STIMULUS	ENVIRONMENT	ARTIFACT	RESPONSE	MEASURE
A delivery engineer	files a ticket that touches five files	a normal run, local models, no cloud	the AIForge agent line	one merged change that builds; the parts link	**0 → 53 linked lines; first full run green**

MORE SCENARIOS

Nothing leaves our network

Check: A check fails the build if it would call out

Running today, small batch

We know the size limit

Check: Files per ticket handled cleanly

3–5 files now

Team memory does not bloat

Check: Duplicate notes; growth per week

Fold step designed

In short: splitting the work beat waiting for a bigger model — and we have the number to show it.

Decision Record (ADR)

How?
Canvas 12 of 14

It answers: what did we decide, what else did we consider, and why?

WHAT IT IS

ADR = architecture decision record. One decision with its options, reasons and cost.

HOW IT HELPS

- Keeps the reasons, so nobody re-argues them
- Forces real options, not one favourite
- Names the bias you checked and the owner

CONTEXT

Our local models hold about 3–5 files at once. A multi-file ticket must still build as one piece — and no code may leave.

REJECTED

Ready-made agent

Broke on our model server; cannot build Java.

REJECTED

One big model

0 linked lines on a multi-file ticket.

REJECTED

Cloud AI

Works — and sends our code outside. Not allowed.

CHOSEN

Small-agent line

0 → 53 linked lines. Builds before every merge.

DECISION — WE WILL…

finish each ticket with a line of small specialist agents, each with one small job. Simple tickets skip the line.

CONSEQUENCES

+ each role is checked on its own; a bad part fails alone
– more model calls, more waiting, more moving parts

REVERSIBILITY



easy to undo — permanent

Bias checked: “more agents must be better”
Owner: project lead

In short: we did not pick it because we liked it — it is the one our tests kept passing.

OKR Card

Where next?
Canvas 13 of 14

It answers: what does good look like, in numbers we can check?

WHAT IT IS

OKR = objective and key results: one goal and a few measures.

HOW IT HELPS

- One goal everyone can repeat
- Measures you can actually check
- Shows honestly where you are today

OBJECTIVE **Make AIForge the normal way small work gets done**

KEY RESULT	PROGRESS TO TARGET	TODAY	TARGET
Tickets that end in a merged pull request		Not measured	3 in 4
Files one ticket can change cleanly		3 to 5	8 to 10
Tickets filed by people other than the leads		Almost none	Half
People who can run and fix the platform		1	2
Runs where data left our network		Zero	Zero

In short: the tool works; now we must measure how often it succeeds. The first row is the problem.

Roadmap Canvas

Where next?
Canvas 14 of 14

It answers: what happens when — for product, platform and people?

WHAT IT IS

The plan over time, with product, platform and people in one view.

HOW IT HELPS

- Change is people and process, not only code
- Each row moves together, quarter by quarter
- Every item traces back to an OKR or a gap

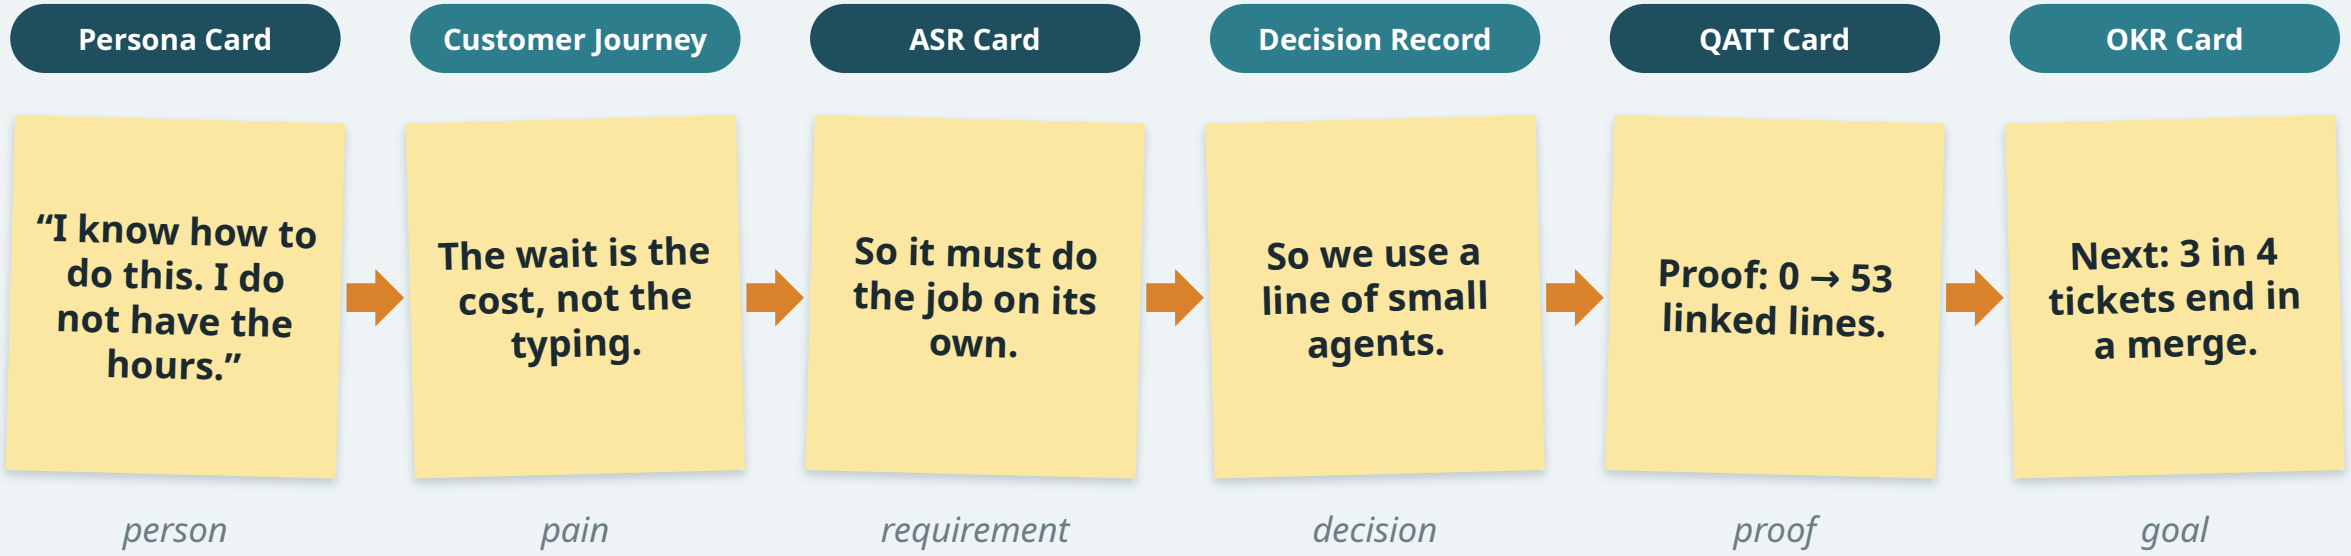
	Q1	Q2	Q3	Q4
Product	Honest verdict on every pull request	Team memory sharing live	Bigger tickets, more files	Open to every team
Platform	GPU upgrade sized	Bigger local models tested	Model service run by IT operations	Capacity planned ahead
People	Second operator named	QA owns the test gate	Juniors onboarded	Adoption measured monthly

In short: one rule for the whole year: nothing on this roadmap moves work off our own hardware.

How the canvases connect

The thread
Cause and effect

One thread through six canvases — each claim has a canvas behind it



Why this matters: change one note and you know which others to revisit. That is the link a long document hides.

In short: the canvases are not fourteen documents. They are one argument on fourteen pages.

What to take away

How to use the canvases on your own work — and what AIForge taught us.

1 Using the canvases

- Start with the person, not the technology
- One page each — if it does not fit, it is not clear yet
- Fill boxes with facts; leave the gaps visible
- Pin them side by side and follow the links

2 What AIForge taught us

- The limit was the design, not the model
- The canvases made the argument checkable — one page, not a long thread
- A tool one person can run is not yet a tool the company owns

Every claim here has a canvas behind it.